

# Machine Learning from Titanic Disaster

Pouria

3/28/2022

## Contents

<b>Objective</b>	<b>2</b>
<b>Import Required Libraries and Tools</b>	<b>3</b>
<b>1. Import/read the data.</b>	<b>4</b>
Import the csv file . . . . .	4
<b>2. Preprocess the data.</b>	<b>4</b>
Missing values . . . . .	4
Leave out irrelevant variables . . . . .	5
Visualize missing values . . . . .	5
Impute the missing values . . . . .	6
<b>3. Split the data</b>	<b>6</b>
Split the data . . . . .	7
<b>4. Logistic Regression</b>	<b>7</b>
Fit and summarize . . . . .	7
Use <code>anova()</code> . . . . .	7
Pseudo-R2 . . . . .	8
<b>5. Evaluate your model on the test set.</b>	<b>8</b>
Misclassification error . . . . .	8
Confusion Matrix . . . . .	9
Plot ROC . . . . .	9
Obtain AUC . . . . .	10
Calculate the accuracy using CV error . . . . .	10

<b>6. Fit a new model</b>	<b>11</b>
Implement a new fit with <b>Pclass</b> and <b>Age</b> . . . . .	11
Anova . . . . .	11
Estimate the accuracy with 10-fold CV. . . . .	11
Misclassification on test set . . . . .	12
Confusion Matrix . . . . .	12
ROC Curve . . . . .	13
Both ROC Curves . . . . .	13

## Objective

This dataset contain the data related to the passengers of the Titanic. The goal is to see if we are able to build a model that can predict whether or not a passenger would survive. The **Survived** column of this dataset reflects if they really survived or not (1: survived, and 0: didn't survive).

1. Load the attached csv file into your Markdown accounting for the missing values.
2. Preprocess the data.
  - Check for the missing values for each variable.
  - Use `missmap()` to visualize the missing values.
  - Reasonably ignore variables that probably do not affect the predictions.
  - Impute the missing values in the numeric (continuous) variables. Only **Age** here.
  - Remove any of the rows with **NA** in their categorical variables (any categorical variable that is left in your data after leaving the unnecessary ones out).
3. Split the data up into the training and test set. (`set.seed = 1`).
  - Sample 100 rows from the data and use as your test set. Use the rest as your training set.
4. Fit a logistic regression model onto the training set
  - Fit and summarize the model to predict **Survived** as a function of all other remaining variables.
  - Use `anova()` to see the effect of each variable on the residual deviance.
  - Obtain pseudo-R2 of the model, e.g., McFadden R2.
5. Evaluate your model on the test set.
  - Calculate the misclassification error on your test set with `threshold = 0.5`
  - the accuracy result your obtained depends on the test set – so, now re-estimate the accuracy with a 10-fold CV
  - Compute the confusion matrix for the model with the test set.
  - Use library `ROCR` to plot ROC curve and also obtain the AUC of your model.
6. Redo the last two steps with only **Pclass** and **Age**.
  - At last, plot the ROC curve from the two models on the same figure. What do you conclude?

## Import Required Libraries and Tools

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(Amelia)
```

```
## Loading required package: Rcpp
```

```
## ##
```

```
## ## Amelia II: Multiple Imputation
```

```
## ## (Version 1.8.1, built: 2022-11-18)
```

```
## ## Copyright (C) 2005-2023 James Honaker, Gary King and Matthew Blackwell
```

```
## ## Refer to http://gking.harvard.edu/amelia/ for more information
```

```
## ##
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
library(boot)
```

```
##
```

```
## Attaching package: 'boot'
```

```
## The following object is masked from 'package:lattice':
```

```
##
```

```
## melanoma
```

```
library(ROCR)
```

```
library(pscl)
```

```
## Classes and Methods for R developed in the
```

```
## Political Science Computational Laboratory
```

```
## Department of Political Science
```

```
## Stanford University
```

```
## Simon Jackman
```

```
## hurdle and zeroinfl functions by Achim Zeileis
```

## 1. Import/read the data.

Load the attached csv file into your Markdown accounting for the missing values.

### Import the csv file

Let's import the training data here:

```
training.data.raw <- read.csv('Titanic.csv', na.strings=c(""))
# training.data.raw <- read.csv('Titanic.csv')

head(training.data.raw)
```

```
##   PassengerId Survived Pclass
## 1           1         0       3
## 2           2         1       1
## 3           3         1       3
## 4           4         1       1
## 5           5         0       3
## 6           6         0       3
##                                     Name    Sex Age SibSp Parch
## 1                                     Braund, Mr. Owen Harris   male  22     1     0
## 2 Cumings, Mrs. John Bradley (Florence Briggs Thayer) female  38     1     0
## 3                                     Heikkinen, Miss. Laina female  26     0     0
## 4 Futrelle, Mrs. Jacques Heath (Lily May Peel) female    35     1     0
## 5                                     Allen, Mr. William Henry   male  35     0     0
## 6                                     Moran, Mr. James         male  NA     0     0
##   Ticket    Fare Cabin Embarked
## 1  A/5 21171  7.2500 <NA>      S
## 2   PC 17599 71.2833   C85      C
## 3 STON/O2. 3101282  7.9250 <NA>      S
## 4   113803 53.1000  C123      S
## 5   373450  8.0500 <NA>      S
## 6   330877  8.4583 <NA>      Q
```

## 2. Preprocess the data.

### Missing values

- Check for the missing values for each variable.

Now we need to check for missing values and look how many unique values there are for each variable using the `sapply()` function to apply `is.na()` on to all variables in the data.

```
# sapply(data, function(x){})
```

```
sapply(training.data.raw,function(x) sum(is.na(x)))
```

```
## PassengerId    Survived    Pclass    Name    Sex    Age
##           0           0           0         0         0    177
##      SibSp      Parch      Ticket    Fare    Cabin Embarked
##           0           0           0         0         687      2
```

```
sapply(training.data.raw, function(x) length(unique(x)))
```

```
## PassengerId    Survived    Pclass    Name    Sex    Age
##      891           2           3      891         2    89
##      SibSp      Parch      Ticket    Fare    Cabin Embarked
##           7           7        681      248      148      4
```

## Leave out irrelevant variables

- Ignore/leave out the variables PassengerID, Cabin, Ticket, and Name.

```
names(training.data.raw)
```

```
## [1] "PassengerId" "Survived"    "Pclass"      "Name"        "Sex"
## [6] "Age"          "SibSp"       "Parch"       "Ticket"      "Fare"
## [11] "Cabin"        "Embarked"
```

```
# data <- subset(training.data.raw, select=c(2,3,5,6,7,8,10,12))
data <- training.data.raw %>%
  select(-c("PassengerId", "Cabin", "Ticket", "Name"))
head(data)
```


```
##   Survived Pclass    Sex Age SibSp Parch   Fare Embarked
## 1         0      3  male  22    1    0  7.2500         S
## 2         1      1 female  38    1    0 71.2833         C
## 3         1      3 female  26    0    0  7.9250         S
## 4         1      1 female  35    1    0 53.1000         S
## 5         0      3  male  35    0    0  8.0500         S
## 6         0      3  male  NA    0    0  8.4583         Q
```

## Visualize missing values

- Use `missmap()` to visualize the missing values.

We can also visualize the missing values using the function `missmap()`.

```
missmap(training.data.raw, main = "Missing values vs observed")
```



## Impute the missing values

- Only keep those rows that are not NA in Embarked
- Then, impute the missing values in Age with mean value.

```
data <- data[!is.na(data$Embarked),]

data$Age[is.na(data$Age)] <- mean(data$Age, na.rm=T)

head(data)
```

```
##   Survived Pclass   Sex    Age SibSp Parch   Fare Embarked
## 1         0      3  male 22.00000     1     0  7.2500         S
## 2         1      1 female 38.00000     1     0 71.2833         C
## 3         1      3 female 26.00000     0     0  7.9250         S
## 4         1      1 female 35.00000     1     0 53.1000         S
## 5         0      3  male 35.00000     0     0  8.0500         S
## 6         0      3  male 29.64209     0     0  8.4583         Q
```

## 3. Split the data

Split the data up into the training and test set. (`set.seed = 1`).

## Split the data

- Sample 100 rows from the data and use as your test set. Use the rest as your training set.

```
set.seed(1)
test.ID <- sample(dim(data)[1], 100)
train <- data[-test.ID,]
test <- data[test.ID,]
```

## 4. Logistic Regression

Fit a logistic regression model onto the training set

### Fit and summarize

- Fit and summarize the model to predict Survived as a function of all other remaining variables.

```
model <- glm(Survived ~ ., family=binomial(link="logit"), data = train)
summary(model)
```

```
##
## Call:
## glm(formula = Survived ~ ., family = binomial(link = "logit"),
##      data = train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  5.246098   0.601126   8.727  < 2e-16 ***
## Pclass       -1.047623   0.152370  -6.876 6.18e-12 ***
## Sexmale      -2.717508   0.212350 -12.797 < 2e-16 ***
## Age          -0.041380   0.008403  -4.925 8.45e-07 ***
## SibSp        -0.298049   0.110674  -2.693  0.00708 **
## Parch        -0.125844   0.124312  -1.012  0.31138
## Fare          0.001772   0.002772   0.639  0.52263
## EmbarkedQ    -0.146313   0.403005  -0.363  0.71656
## EmbarkedS    -0.428994   0.247267  -1.735  0.08275 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1049.04  on 788  degrees of freedom
## Residual deviance:  706.49  on 780  degrees of freedom
## AIC: 724.49
##
## Number of Fisher Scoring iterations: 5
```

### Use anova()

- Use `anova()` to see the effect of each variable on the residual deviance.

```
anova(model, test="Chisq")
```

```
## Analysis of Deviance Table
##
## Model: binomial, link: logit
##
## Response: Survived
##
## Terms added sequentially (first to last)
##
##
##          Df Deviance Resid. Df Resid. Dev  Pr(>Chi)
## NULL                788      1049.04
## Pclass    1    82.196      787      966.85 < 2.2e-16 ***
## Sex       1   221.992      786      744.85 < 2.2e-16 ***
## Age       1    20.725      785      724.13 5.302e-06 ***
## SibSp     1    12.431      784      711.70 0.0004223 ***
## Parch     1     0.857      783      710.84 0.3544538
## Fare      1     0.974      782      709.87 0.3236963
## Embarked  2     3.376      780      706.49 0.1848942
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Pseudo-R2

- Obtain pseudo-R2 of the model, e.g., McFadden R2.

```
pR2(model)
```

```
## fitting null model for pseudo-r2
##
##          llh          llhNull          G2          McFadden          r2ML          r2CU
## -353.2453603 -524.5205194  342.5503183    0.3265366    0.3521898    0.4789011
```

## 5. Evaluate your model on the test set.

### Misclassification error

- Use your model on the test set to make predictions of Survived.
- Calculate the misclassification error on your test set with threshold = 0.5.

```
pred.results <- predict(model, newdata=subset(test, select=-1), type = "response")
pred.results <- ifelse(pred.results > .5, 1, 0)

misClassificationError <- mean(pred.results!= test$Survived)
print(paste("Accuracy = ", 1-misClassificationError))

## [1] "Accuracy = 0.83"
```



## Confusion Matrix

- Compute the confusion matrix for the model with the test set.

You may use either `table()` or `confusionMatrix()` to obtain the confusion matrix.

```
confusionMatrix(factor(test$Survived), factor(pred.results))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 52  9
##           1  8 31
##
##           Accuracy : 0.83
##           95% CI : (0.7418, 0.8977)
##    No Information Rate : 0.6
##    P-Value [Acc > NIR] : 6.256e-07
##
##           Kappa : 0.6444
##
##  McNemar's Test P-Value : 1
##
##           Sensitivity : 0.8667
##           Specificity : 0.7750
##           Pos Pred Value : 0.8525
##           Neg Pred Value : 0.7949
##           Prevalence : 0.6000
##           Detection Rate : 0.5200
##    Detection Prevalence : 0.6100
##           Balanced Accuracy : 0.8208
##
##           'Positive' Class : 0
##
```


```
# table(test$Survived, pred.results)
```

## Plot ROC

- Use library **ROCR** to plot ROC curve.

Use `prediction()` and then pass its output to `performance()` to obtain the true positive rate ("**tpr**") and false positive rate ("**fpr**") values for different thresholds. Notice that you will not have to worry about setting thresholds. `performance()` will automatically try different thresholds and accordingly return the rate values in two arrays that are accessible by `@` (instead of `$`).

```
pred <- predict(model, newdata=subset(test, select=-1), type = "response")
pr <- prediction(pred, test$Survived)
prf <- performance(pr, measure = "tpr", x.measure = "fpr")
plot(prf)
```



## Obtain AUC

- Your measure in `performance()` will be "auc" this time.

```
auc <- performance(pr, measure = "auc")
auc <- auc@y.values[[1]]
auc
```

```
## [1] 0.9058428
```

## Calculate the accuracy using CV error

- the accuracy result your obtained depends on the test set – so, now re-estimate the accuracy with a 10-fold CV

```
model.full <- glm(Survived ~ ., family=binomial(link="logit"), data = data)
cv.error <- cv.glm(data, model.full, K=10)$delta[1]
accuracy <- 1 - cv.error
accuracy
```

```
## [1] 0.8566667
```

## 6. Fit a new model

- Redo the last two steps with only Pclass and Age.

### Implement a new fit with Pclass and Age

Do the fit again with only Pclass and Age;

```
model2 <- glm(Survived ~ Pclass + Age, family=binomial(link="logit"), data = train)
summary(model2)
```

```
##
## Call:
## glm(formula = Survived ~ Pclass + Age, family = binomial(link = "logit"),
##      data = train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  2.92568    0.37483   7.805 5.94e-15 ***
## Pclass      -1.03179    0.10490  -9.836 < 2e-16 ***
## Age         -0.03666    0.00678  -5.408 6.38e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 1049.04  on 788  degrees of freedom
## Residual deviance:  935.76  on 786  degrees of freedom
## AIC: 941.76
##
## Number of Fisher Scoring iterations: 4
```

### Anova

```
anova(model, model2)
```

```
## Analysis of Deviance Table
##
## Model 1: Survived ~ Pclass + Sex + Age + SibSp + Parch + Fare + Embarked
## Model 2: Survived ~ Pclass + Age
##   Resid. Df Resid. Dev Df Deviance
## 1         780      706.49
## 2         786      935.76 -6  -229.27
```

Estimate the accuracy with 10-fold CV.

```
model2.cv <- glm(Survived ~ Pclass + Age, family=binomial(link="logit"), data = data)
cv.error2 <- cv.glm(data, model2.cv, K=10)$delta[1]
accuracy <- 1 - cv.error2
accuracy
```

```
## [1] 0.7985145
```

## Misclassification on test set

- Make predictions on the test set and calculate misclassification

```
pred.results2 <- predict(model2, newdata=subset(test, select=-1))
pred.results2 <- ifelse(pred.results > .5, 1, 0)

misClassificationError <- mean(pred.results != test$Survived)
print(paste("Accuracy = ", 1-misClassificationError))
```

```
## [1] "Accuracy = 0.83"
```

## Confusion Matrix


```
confusionMatrix(factor(test$Survived), factor(pred.results2))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0  1
##           0 52  9
##           1  8 31
##
##           Accuracy : 0.83
##           95% CI : (0.7418, 0.8977)
##           No Information Rate : 0.6
##           P-Value [Acc > NIR] : 6.256e-07
##
##           Kappa : 0.6444
##
##           McNemar's Test P-Value : 1
##
##           Sensitivity : 0.8667
##           Specificity : 0.7750
##           Pos Pred Value : 0.8525
##           Neg Pred Value : 0.7949
##           Prevalence : 0.6000
##           Detection Rate : 0.5200
##           Detection Prevalence : 0.6100
##           Balanced Accuracy : 0.8208
##
##           'Positive' Class : 0
##
```

## ROC Curve

Plot ROC for this new model on top of the ROC curve for the full model - what do you conclude?

```
pred2 <- predict(model2, newdata=subset(test, select=-1), type = "response")
pr2 <- prediction(pred2, test$Survived)
prf2 <- performance(pr2, measure = "tpr", x.measure = "fpr")
plot(prf2)
```



```
auc2 <- performance(pr2, measure = "auc")
auc2 <- auc2@y.values[[1]]
auc2
```

```
## [1] 0.7881463
```

## Both ROC Curves

```
df1 <- data.frame(prf@y.values, prf@x.values) %>%
  mutate(Model = "Survived ~ .") %>%
  `colnames<-`(c("TP", "FP", "Model"))

df2 <- data.frame(prf2@y.values, prf2@x.values) %>%
  mutate("Model" = "Survived ~ Pclass + Age") %>%
```

```
`colnames<-`(c("TP", "FP", "Model"))

df <- rbind(df1, df2)
```

```
ggplot(df) +
  geom_line(aes(x = FP, y = TP, col = Model), size = 1.5) +
  scale_color_manual(values = c("red", "black"))
```

```
## Warning: Using 'size' aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use 'linewidth' instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```

